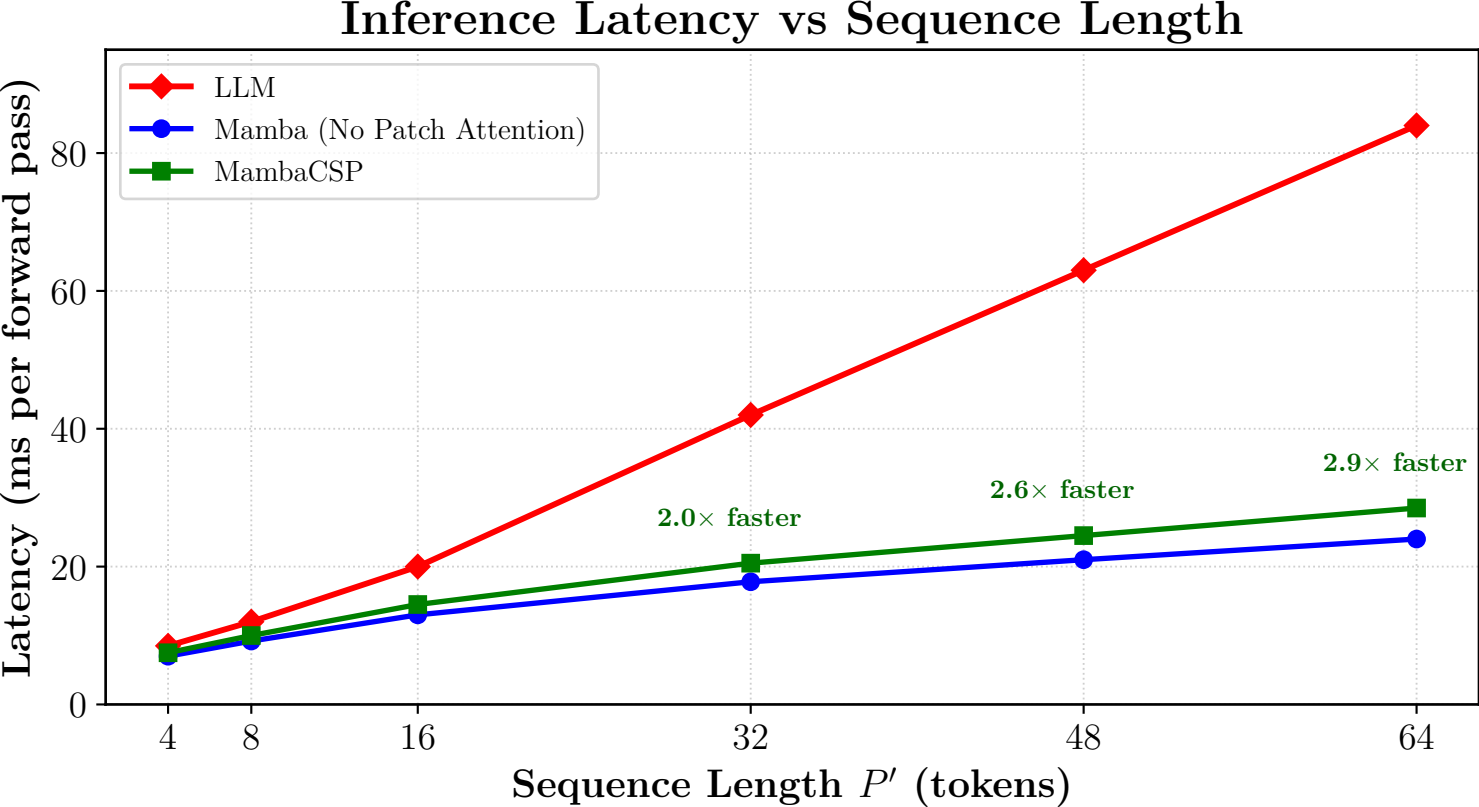
Inference Latency vs Sequence Length
LLM
Mamba (No Patch Attention)
MambaCSP
2.0× faster
2.6× faster
2.9× faster
Latency (ms per forward pass)
0
20
40
60
80
4
8
16
32
48
64
Sequence Length P' (tokens)