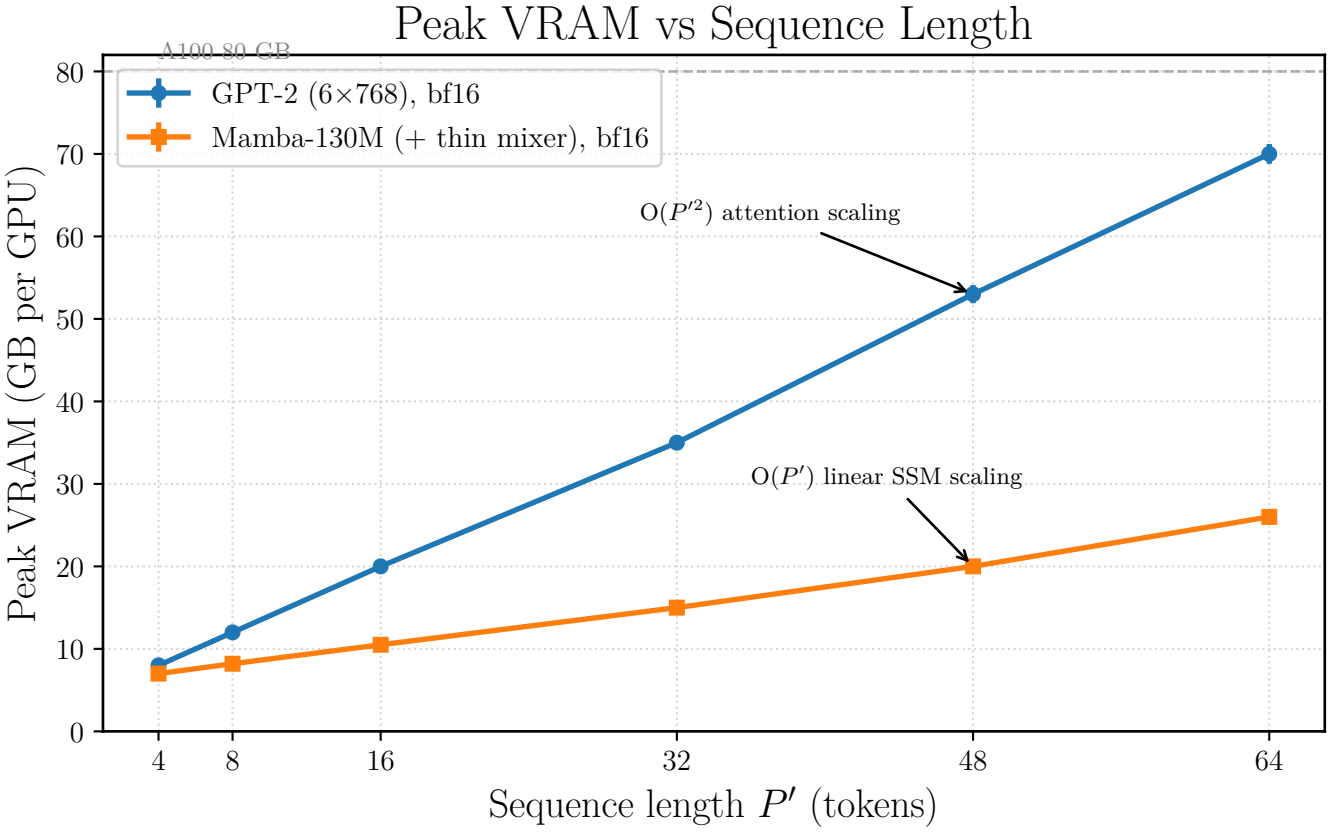
Peak VRAM vs Sequence Length
A100-80 GB
GPT-2 (6×768), bf16
Mamba-130M (+ thin mixer), bf16
O(P'^2) attention scaling
O(P') linear SSM scaling
Peak VRAM (GB per GPU)
0
10
20
30
40
50
60
70
80
4
8
16
32
48
64
Sequence length P' (tokens)